

PCOS DETECTION USING ULTRASOUND IMAGES

1. ResNet-50

- **Architecture:** ResNet-50 is a deep convolutional neural network (CNN) with 50 layers, designed using a **residual learning framework**. The core idea is the use of **residual connections** (also known as skip connections) that allow gradients to flow more easily through the network during backpropagation, solving the vanishing gradient problem and enabling the training of much deeper networks.
- **Key Components:**
 - **Convolutional layers:** Multiple convolutional layers are stacked to learn hierarchical features.
 - **Residual Blocks:** Introduces skip connections that bypass some layers, enabling the network to learn residual mappings instead of full mappings, making it easier to train deeper networks.
 - **Global Average Pooling:** Reduces the dimensionality of the feature maps before fully connected layers.

```
61/61 ————— 11s 113ms/step - accuracy: 0.9869 - loss: 0.0277
Test Loss: 0.02654281258583069
Test Accuracy: 0.9916753172874451
61/61 ————— 11s 133ms/step
```

2. CNN with Pre-Trained VGGNet16

- **Architecture:** VGG16 is a CNN architecture with **16 layers**, including 13 convolutional layers and 3 fully connected layers. It uses small **3x3 convolution filters** stacked on top of each other with max-pooling layers interleaved.
- **Key Components:**
 - **Convolutional layers:** Use a simple and uniform architecture with small filters (3x3), which helps in capturing more fine-grained features.
 - **Pre-trained Weights:** VGGNet16 can be used as a **pre-trained model**, where the model is initially trained on large datasets like ImageNet, and the learned features are transferred to other tasks, such as medical image classification.

- **Global Average Pooling:** Sometimes added on top of the convolutional layers to reduce the dimensionality of the final feature maps.

```
Validation Accuracy: 99.65398%
Validation F1 Score: 0.99639
Validation Precision: 0.99569
Validation Recall: 0.99713
```

```
Model loaded successfully
61/61 ————— 10s 126ms/step
Test Accuracy: 99.17%
```

3. Stacked Ensemble with XGBoost Meta-Learner

1. Architecture and Approach:

- **Stacked Ensemble Architecture:** The model combines three base learners (Logistic Regression, Random Forest, SVM), and their probabilistic predictions are stacked (concatenated). These stacked predictions are then fed into a meta-learner (XGBoost) that makes the final prediction.
- **Approach:** Each base model learns from the data independently, and the meta-learner aggregates their predictions to improve accuracy, leveraging diverse learning techniques (linear, tree-based, and margin-based models).

2. Key Components:

- **Base Learners:** Logistic Regression (linear), Random Forest (non-linear, tree-based), and SVM (margin-based) are used to capture different patterns in the data.
- **Meta-Learner:** XGBoost (gradient-boosted decision trees) integrates the predictions of the base learners to enhance the overall predictive performance.

```
[395] validation_0-logloss:0.39828 validation_0-error:0.00000 validation_1-logloss:0.40684 validation_1-error:0.01730
[396] validation_0-logloss:0.39779 validation_0-error:0.00000 validation_1-logloss:0.40638 validation_1-error:0.01730
[397] validation_0-logloss:0.39730 validation_0-error:0.00000 validation_1-logloss:0.40591 validation_1-error:0.01730
[398] validation_0-logloss:0.39681 validation_0-error:0.00000 validation_1-logloss:0.40545 validation_1-error:0.01730
[399] validation_0-logloss:0.39632 validation_0-error:0.00000 validation_1-logloss:0.40498 validation_1-error:0.01730
Accuracy: 0.9973985431841832
F1 Score: 0.997306658865263
Precision: 0.9968193384223918
Recall: 0.9978089395267309
ROC AUC: 0.997808939526731
```

+ Code + Markdown

4. Multi-Model Ensemble (VGG16 + ResNet50 + RF + SVM + Logistic Regression + XGBoost)

1. Architecture and Approach:

This model implements a stacking-based ensemble for binary classification using ultrasound images for PCOS detection. **Features are extracted** from two pre-trained deep learning architectures, **VGG16 and ResNet50**, which capture high-level representations from the images. These extracted features are then combined and fed into three different base classifiers: **Logistic Regression (LR), Random Forest (RF), and Support Vector Machine (SVM)**. The outputs of these base models (probabilities) are stacked together and passed to a **meta-learner (XGBoost) for final classification**, allowing for improved predictive performance.

2. Key Components:

- **Feature Extractors:** VGG16 and ResNet50 are used to extract features from images, leveraging their deep learning architectures pre-trained on ImageNet.
- **Base Classifiers:** Logistic Regression, Random Forest, and Support Vector Machine are used as the initial classifiers in the stacking ensemble, providing multiple perspectives on the input data.
- **Meta-Learner:** XGBoost, a powerful gradient boosting algorithm, serves as the final classifier in the ensemble, combining the predictions from the base models to make more accurate final predictions.

Output 1:

```
lr = LogisticRegression(max_iter=500)
rf = RandomForestClassifier(n_estimators=100)
svm = SVC(probability=True)
meta_learner = xgb.XGBClassifier(n_estimators=200, learning_rate=0.01,
max_depth=4, eval_metric=['logloss', 'error'])
```

```
[195] validation_0-logloss:0.07108 validation_0-error:0.00000 validation_1-logloss:0.13811 validation_1-error:0.02708
[196] validation_0-logloss:0.07095 validation_0-error:0.00000 validation_1-logloss:0.13768 validation_1-error:0.02768
[197] validation_0-logloss:0.07023 validation_0-error:0.00000 validation_1-logloss:0.13725 validation_1-error:0.02768
[198] validation_0-logloss:0.06952 validation_0-error:0.00000 validation_1-logloss:0.13683 validation_1-error:0.02768
[199] validation_0-logloss:0.06882 validation_0-error:0.00000 validation_1-logloss:0.13642 validation_1-error:0.02768
Test Accuracy: 99.58377%
Test F1 Score: 0.99569
Test Precision: 0.99493
Test Recall: 0.99649
Test ROC AUC: 0.99649
```

[+ Code](#)[+ Markdown](#)**Output 2:**

meta_learner = xgb.XGBClassifier(n_estimators=200, learning_rate=0.01, max_depth=6, eval_metric=['logloss', 'error'])

```
[195] validation_0-logloss:0.07108 validation_0-error:0.00000 validation_1-logloss:0.13479 validation_1-error:0.03460
[196] validation_0-logloss:0.07095 validation_0-error:0.00000 validation_1-logloss:0.13442 validation_1-error:0.03460
[197] validation_0-logloss:0.07023 validation_0-error:0.00000 validation_1-logloss:0.13407 validation_1-error:0.03460
[198] validation_0-logloss:0.06952 validation_0-error:0.00000 validation_1-logloss:0.13372 validation_1-error:0.03460
[199] validation_0-logloss:0.06882 validation_0-error:0.00000 validation_1-logloss:0.13338 validation_1-error:0.03460
Test Accuracy: 99.47971%
Test F1 Score: 0.99462
Test Precision: 0.99368
Test Recall: 0.99562
Test ROC AUC: 0.99562
```

[+ Code](#)[+ Markdown](#)**Output 3:**

meta_learner = xgb.XGBClassifier(n_estimators=500, learning_rate=0.01, max_depth=4, eval_metric=['logloss', 'error'])

```
[495] validation_0-logloss:0.00447 validation_0-error:0.00000 validation_1-logloss:0.13055 validation_1-error:0.02422
[496] validation_0-logloss:0.00443 validation_0-error:0.00000 validation_1-logloss:0.13070 validation_1-error:0.02422
[497] validation_0-logloss:0.00440 validation_0-error:0.00000 validation_1-logloss:0.13086 validation_1-error:0.02422
[498] validation_0-logloss:0.00436 validation_0-error:0.00000 validation_1-logloss:0.13102 validation_1-error:0.02422
[499] validation_0-logloss:0.00433 validation_0-error:0.00000 validation_1-logloss:0.13117 validation_1-error:0.02422
Test Accuracy: 99.63580%
Test F1 Score: 0.99623
Test Precision: 0.99556
Test Recall: 0.99693
Test ROC AUC: 0.99693
```

[+ Code](#)[+ Markdown](#)**Output 4:**

meta_learner = xgb.XGBClassifier(n_estimators=200, learning_rate=0.05, max_depth=4, eval_metric=['logloss', 'error'])

```
[195] validation_0-logloss:0.00122 validation_0-error:0.00000 validation_1-logloss:0.22616 validation_1-error:0.03460
[196] validation_0-logloss:0.00122 validation_0-error:0.00000 validation_1-logloss:0.22616 validation_1-error:0.03460
[197] validation_0-logloss:0.00122 validation_0-error:0.00000 validation_1-logloss:0.22616 validation_1-error:0.03460
[198] validation_0-logloss:0.00122 validation_0-error:0.00000 validation_1-logloss:0.22616 validation_1-error:0.03460
[199] validation_0-logloss:0.00122 validation_0-error:0.00000 validation_1-logloss:0.22616 validation_1-error:0.03460
Test Accuracy: 99.47971%
Test F1 Score: 0.99462
Test Precision: 0.99368
Test Recall: 0.99562
Test ROC AUC: 0.99562
```

[+ Code](#)[+ Markdown](#)

Output 5:

meta_learner = xgb.XGBClassifier(n_estimators=200, learning_rate=0.001, max_depth=4, eval_metric=['logloss', 'error'])

```
[195] validation_0-logloss:0.51441 validation_0-error:0.00000 validation_1-logloss:0.51735 validation_1-error:0.01730
[196] validation_0-logloss:0.51373 validation_0-error:0.00000 validation_1-logloss:0.51670 validation_1-error:0.01730
[197] validation_0-logloss:0.51305 validation_0-error:0.00000 validation_1-logloss:0.51605 validation_1-error:0.01730
[198] validation_0-logloss:0.51237 validation_0-error:0.00000 validation_1-logloss:0.51540 validation_1-error:0.01730
[199] validation_0-logloss:0.51169 validation_0-error:0.00000 validation_1-logloss:0.51475 validation_1-error:0.01730
Test Accuracy: 99.73985%
Test F1 Score: 0.99731
Test Precision: 0.99682
Test Recall: 0.99781
Test ROC AUC: 0.99781
```


+ Code

+ Markdown

Output 5:

meta_learner = xgb.XGBClassifier(n_estimators=500, learning_rate=0.001, max_depth=4, eval_metric=['logloss', 'error'])

```
[497] validation_0-logloss:0.35107 validation_0-error:0.00000 validation_1-logloss:0.35323 validation_1-error:0.01730
[498] validation_0-logloss:0.35124 validation_0-error:0.00000 validation_1-logloss:0.35482 validation_1-error:0.01730
[499] validation_0-logloss:0.35082 validation_0-error:0.00000 validation_1-logloss:0.35441 validation_1-error:0.01730
Test Accuracy: 99.89594%
Test F1 Score: 0.99892
Test Precision: 0.99872
Test Recall: 0.99912
Test ROC AUC: 0.99912
```

Load and Preprocess Images

|

Evaluate Model